

Introduction to Big Data and Data Science (CSCE 5300)*

Yunhe Feng

Assistant Professor, Department of Computer Science and Engineering

12th September, 2024



- 1 Motivation of DataFrame, Pandas, and PySpark
- 2 Pandas and DataFrame
- 3 PySpark
- 4 Assignment

How to Process a CSV File with 100,000 Rows



If someone gives you a CSV file with 100,000 rows in it, what tools do you use to start exploring and understanding that data?

1:03 PM · Sep 20, 2022 · Twitter for iPhone

907 Retweets 679 Quote Tweets 7,726 Likes

How to Process a CSV File with 100,000 Rows

 **Simon Willison** ✓ @simonw · Sep 20 ...
If someone gives you a CSV file with 100,000 rows in it, what tools do you use to start exploring and understanding that data?

💬 2,712 ↺ 1,585 ❤️ 7,721 📎

[Show this thread](#)

 **Simon Willison** ✓ @simonw · Sep 20 ...
My own answer: I either open the CSV directly in the **Datasette** Desktop Mac application (datasette.io/desktop) or I do this:

```
sqlite-utils insert /tmp/data.db rows big.csv --csv
datasette /tmp/data.db
```

That gives me a table called "rows" in a fresh SQLite database

💬 14 ↺ 47 ❤️ 869 📎

 **Simon Willison** ✓ @simonw · Sep 20 ...
I tend to start this kind of thing in /tmp because otherwise I end up with hundreds of temporary database files littering my file system - if the exploration starts turning into a project I move it to a more permanent location (backed up by Dropbox)

💬 19 ↺ 3 ❤️ 195 📎

How to Process a CSV File with 100,000 Rows



Matt Hodges @hodgesmr · Sep 20

Replying to @simonw

In order:

1. wc -l
2. head
3. less
4. VisiData
5. Google Sheets (optional)
6. **Pandas**

4



140



...



Jeremy Howard @jeremyphoward · Sep 20

Replying to @simonw

I use **pandas** and sklearn random forests

8



4

130



Ron Itelman @ron_itelman · Sep 20

@jeremyphoward if you could create a simple example of your process, I'd love to learn!

2



6



Rajko Radovanović @rajko_rad · Sep 20

Replying to @simonw

Excel... as would 99% of the world, don't think This thread is super representative 😊

Once you hit 1M, I'd usually use **pandas**, but less technical folks I've seen use Microsoft Access... also if I wanted to plot things locally I would go to tools like Tableau public version...

8



5

191



What is pandas



What is pandas



pandas is a fast, powerful, flexible and easy to use open source data analysis and manipulation tool, built on top of the Python programming language.

pandas VS NumPy



- pandas is built on the NumPy library and written in languages like Python, Cython, and C. It is a convenience wrapper over NumPy.
- NumPy is a Python library used for working with arrays.

Why is pandas Called pandas

- **pandas**: The library's name derives from **panel data**, a common term for multidimensional data sets encountered in statistics and econometrics¹.
- **Panel Data**: In statistics and econometrics, panel data is multi-dimensional data involving measurements over time. Panel data is a subset of longitudinal data where observations are for the same subjects each time. – Wikipedia

person ↕	year ↕	income ↕	age ↕	sex ↕
1	2016	1300	27	1
1	2017	1600	28	1
1	2018	2000	29	1
2	2016	2000	38	2
2	2017	2300	39	2
2	2018	2400	40	2

¹“pandas: a Foundational Python Library for Data Analysis and Statistics” by Wes McKinney

pandas VS PySpark



- pandas run operations on a single machine whereas PySpark runs on multiple machines.
- The **DataFrame** structure in PySpark is conceptually similar to (and inspired by) the one in pandas.

- 1 Motivation of DataFrame, Pandas, and PySpark
- 2 Pandas and DataFrame
- 3 PySpark
- 4 Assignment

pandas Data Structures

- Series: a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the index.
- DataFrame: a 2-dimensional labeled data structure with columns of potentially different types.

Series: From ndarray

```
In [3]: s = pd.Series(np.random.randn(5), index=["a", "b", "c", "d", "e"])
```

```
In [4]: s
```

```
Out[4]:
```

```
a    0.469112  
b   -0.282863  
c   -1.509059  
d   -1.135632  
e    1.212112  
dtype: float64
```

```
In [5]: s.index
```

```
Out[5]: Index(['a', 'b', 'c', 'd', 'e'], dtype='object')
```

```
In [6]: pd.Series(np.random.randn(5))
```

```
Out[6]:
```

```
0   -0.173215  
1    0.119209  
2   -1.044236  
3   -0.861849  
4   -2.104569  
dtype: float64
```

Series: From dict

```
In [7]: d = {"b": 1, "a": 0, "c": 2}
```

```
In [8]: pd.Series(d)
```

```
Out[8]:
```

```
b    1  
a    0  
c    2  
dtype: int64
```

Series: From dict

```
In [9]: d = {"a": 0.0, "b": 1.0, "c": 2.0}
```

```
In [10]: pd.Series(d)
```

```
Out[10]:
```

```
a    0.0
```

```
b    1.0
```

```
c    2.0
```

```
dtype: float64
```

```
In [11]: pd.Series(d, index=["b", "c", "d", "a"])
```

```
Out[11]:
```

```
b    1.0
```

```
c    2.0
```

```
d    NaN
```

```
a    0.0
```

```
dtype: float64
```

Series: From scalar value

```
In [12]: pd.Series(5.0, index=["a", "b", "c", "d", "e"])
```

```
Out[12]:
```

```
a    5.0
```

```
b    5.0
```

```
c    5.0
```

```
d    5.0
```

```
e    5.0
```

```
dtype: float64
```

>>>



Series is ndarray-like

```
In [13]: s[0]
Out[13]: 0.4691122999071863
```

```
In [14]: s[:3]
Out[14]:
a    0.469112
b   -0.282863
c   -1.509059
dtype: float64
```

```
In [15]: s[s > s.median()]
Out[15]:
a    0.469112
e    1.212112
dtype: float64
```

```
In [16]: s[[4, 3, 1]]
Out[16]:
e    1.212112
d   -1.135632
b   -0.282863
dtype: float64
```

```
In [17]: np.exp(s)
Out[17]:
a    1.598575
b    0.753623
c    0.221118
d    0.321219
e    3.360575
dtype: float64
```

Series is dict-like

```
In [21]: s["a"]
Out[21]: 0.4691122999071863

In [22]: s["e"] = 12.0

In [23]: s
Out[23]:
a      0.469112
b     -0.282863
c     -1.509059
d     -1.135632
e     12.000000
dtype: float64

In [24]: "e" in s
Out[24]: True

In [25]: "f" in s
Out[25]: False
```

DataFrame: From dict of Series or dicts

```
In [38]: d = {
...:     "one": pd.Series([1.0, 2.0, 3.0], index=["a", "b", "c"]),
...:     "two": pd.Series([1.0, 2.0, 3.0, 4.0], index=["a", "b", "c", "d"]),
...: }
...:

In [39]: df = pd.DataFrame(d)

In [40]: df
Out[40]:
   one  two
a  1.0  1.0
b  2.0  2.0
c  3.0  3.0
d  NaN  4.0

In [41]: pd.DataFrame(d, index=["d", "b", "a"])
Out[41]:
   one  two
d  NaN  4.0
b  2.0  2.0
a  1.0  1.0

In [42]: pd.DataFrame(d, index=["d", "b", "a"], columns=["two", "three"])
Out[42]:
   two three
d  4.0  NaN
b  2.0  NaN
a  1.0  NaN
```



DataFrame: From dict of ndarrays / lists

```
In [45]: d = {"one": [1.0, 2.0, 3.0, 4.0], "two": [4.0, 3.0, 2.0, 1.0]}
```

```
In [46]: pd.DataFrame(d)
```

```
Out[46]:
```

	one	two
0	1.0	4.0
1	2.0	3.0
2	3.0	2.0
3	4.0	1.0

```
In [47]: pd.DataFrame(d, index=["a", "b", "c", "d"])
```

```
Out[47]:
```

	one	two
a	1.0	4.0
b	2.0	3.0
c	3.0	2.0
d	4.0	1.0

DataFrame: Column Selection

```
In [72]: df["one"]
Out[72]:
a    1.0
b    2.0
c    3.0
d    NaN
Name: one, dtype: float64

In [73]: df["three"] = df["one"] * df["two"]

In [74]: df["flag"] = df["one"] > 2

In [75]: df
Out[75]:
```

	one	two	three	flag
a	1.0	1.0	1.0	False
b	2.0	2.0	4.0	False
c	3.0	3.0	9.0	True
d	NaN	4.0	NaN	False

DataFrame: Column Deletion

```
In [76]: del df["two"]
```

```
In [77]: three = df.pop("three")
```

```
In [78]: df
```

```
Out[78]:
```

	one	flag
a	1.0	False
b	2.0	False
c	3.0	True
d	NaN	False

DataFrame: Column Addition

```
In [79]: df["foo"] = "bar"
```

```
In [80]: df
```

```
Out[80]:
```

	one	flag	foo
a	1.0	False	bar
b	2.0	False	bar
c	3.0	True	bar
d	NaN	False	bar

DataFrame: Indexing / Selection

Operation	Syntax	Result
Select column	<code>df[col]</code>	Series
Select row by label	<code>df.loc[label]</code>	Series
Select row by integer location	<code>df.iloc[loc]</code>	Series
Slice rows	<code>df[5:10]</code>	DataFrame
Select rows by boolean vector	<code>df[bool_vec]</code>	DataFrame

- ◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ≡ ≡ ↺ 🔍 ↻

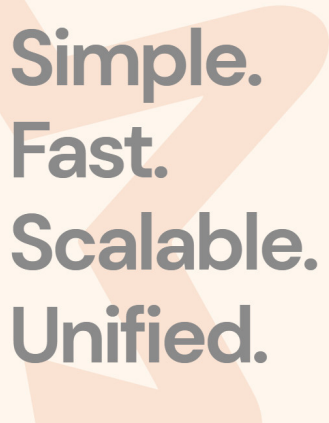
What is PySpark




What is Spark

Apache Spark is a data processing framework that can **quickly** perform processing tasks on very **large** data sets, and can also **distribute** data processing tasks across **multiple computers**, either on its own or in tandem with other distributed computing tools.

What is Spark




Simple.
Fast.
Scalable.
Unified.




Batch/streaming data

Unify the processing of your data in batches and real-time streaming, using your preferred language:
Python, SQL, Scala, Java or R.




SQL analytics

Execute fast, distributed ANSI SQL queries for dashboarding and ad-hoc reporting. Runs faster than most data warehouses.



Data science at scale

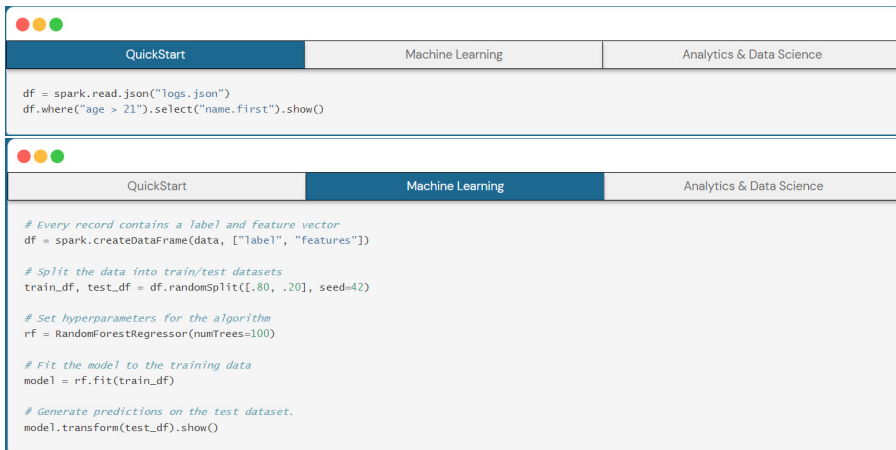
Perform Exploratory Data Analysis (EDA) on petabyte-scale data without having to resort to downsampling



Machine learning

Train machine learning algorithms on a laptop and use the same code to scale to fault-tolerant clusters of thousands of machines.

What is Spark



The image displays two screenshots of a PySpark IDE interface. The top screenshot shows the 'QuickStart' tab selected, with a code editor containing the following code:

```
df = spark.read.json("logs.json")
df.where("age > 21").select("name.first").show()
```

The bottom screenshot shows the 'Machine Learning' tab selected, with a code editor containing the following code:

```
# Every record contains a label and feature vector
df = spark.createDataFrame(data, ["label", "features"])


# Split the data into train/test datasets
train_df, test_df = df.randomSplit([.80, .20], seed=42)

# Set hyperparameters for the algorithm
rf = RandomForestRegressor(numTrees=100)

# Fit the model to the training data
model = rf.fit(train_df)

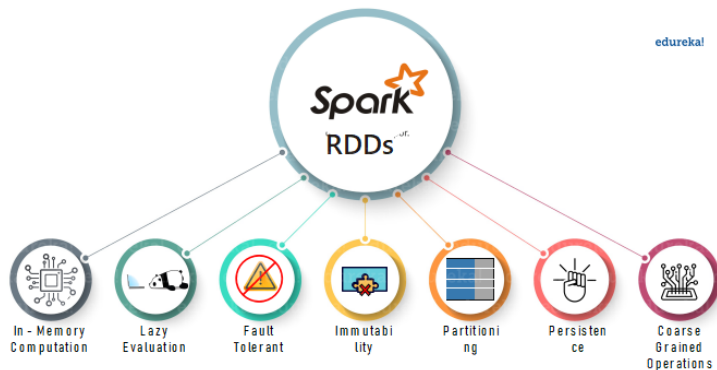
# Generate predictions on the test dataset.
model.transform(test_df).show()
```

Spark Architecture



- Driver Program: create SparkContext and translate the user-written code into jobs
- Cluster Manager: allocate resource and split a job into multiple smaller tasks
- Worker Nodes: execute the tasks

Spark Resilient Distributed Dataset (RDD)



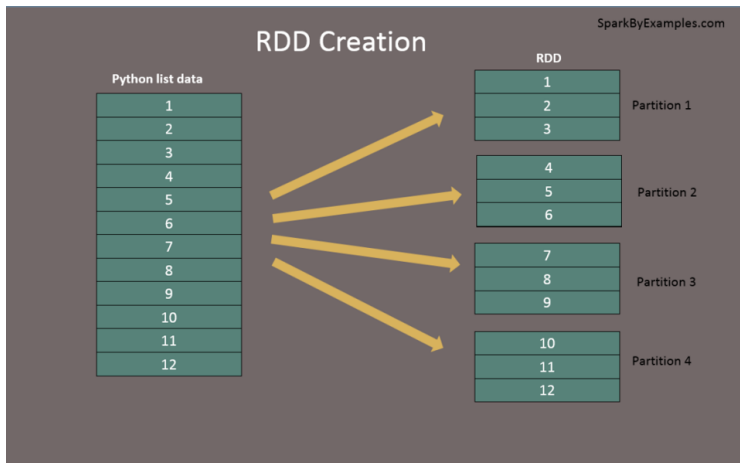
- RDD: a programming abstraction that represents an immutable collection of objects that can be split across a computing cluster
- Operations on RDDs: can also be split across the cluster and executed in a parallel batch process

RDD: Parallelize()

- Parallelize() is a function in SparkContext
- Create Resilient Distributed Datasets (RDD) from a list collection
- Each dataset in RDD is divided into logical partitions, each partition may be computed on a different node of the cluster

```
from pyspark import SparkContext, SparkConf
nums = list(range(0,100000,1))
sc = sparkContext.parallelize(nums)
```


RDD: Parallelize()



RDD VS DataFrame

Feature	RDD	DataFrame	Dataset
Immutable	Yes	Yes	Yes
Fault Tolerant	Yes	Yes	Yes
Type-Safe	Yes	No	Yes
Schema	No	Yes	Yes
Execution Optimization	No	Yes	Yes
Optimizer Engine	N/A	Catalyst Engine	Catalyst Engine
API Level for manipulating distributed collection of data	Low	High	High
language Support	Java, Scala, Pyt	Java, Scala, Python, R	Java, Scala

Install PySpark on Linux

- <https://sparkbyexamples.com/pyspark/install-pyspark-in-anaconda-jupyter-notebook/>
- https://spark.apache.org/docs/latest/api/python/getting_started/install.html
- <https://zhangdijohn.medium.com/pyspark-3-ubuntu-20-04-installation-2792c5c221de>

Install PySpark on Linux

```
yunhe@yunhe-desktop: ~  
(base) yunhe@yunhe-desktop:~$ pyspark  
Python 3.8.8 (default, Apr 13 2021, 19:58:26)  
[GCC 7.3.0] :: Anaconda, Inc. on linux  
Type "help", "copyright", "credits" or "license" for more information.  
22/09/22 21:12:24 WARN Utils: Your hostname, yunhe-desktop resolves to a loopback  
address: 127.0.1.1; using 130.45.5.27 instead (on interface enp7s0)  
22/09/22 21:12:24 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another  
address  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel  
(newLevel).  
22/09/22 21:12:25 WARN NativeCodeLoader: Unable to load native-hadoop library fo  
r your platform... using builtin-java classes where applicable  
Welcome to  
  
      _ _ _ _ _  
     / _ _ _ _ \   version 3.3.0  
    / _ _ _ _ \  
   / _ _ _ _ \  
  / _ _ _ _ \  
 / _ _ _ _ \  
/_ _ _ _ _  
Using Python version 3.8.8 (default, Apr 13 2021 19:58:26)  
Spark context Web UI available at http://130-45-5-27.dyn.grandenetworks.net:4040  
Spark context available as 'sc' (master = local[*], app id = local-1663906346020)
```

PySpark Coding Examples

```
import pyspark
sc = pyspark.SparkContext('local[*]')
big_list = range(10000)
rdd = sc.parallelize(big_list, 2)
odds = rdd.filter(lambda x: x % 2 != 0)
output = odds.take(5)
print(output)
```

```
(base) yunhe@yunhe-desktop:~$ vim test_2.py
(base) yunhe@yunhe-desktop:~$ ~/anaconda3/bin/python3.8 test_2.py
22/09/22 21:32:40 WARN Utils: Your hostname, yunhe-desktop resolves to a loopbac
k address: 127.0.1.1; using 130.45.5.27 instead (on interface enp7s0)
22/09/22 21:32:40 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another
address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel
l(newLevel).
22/09/22 21:32:41 WARN NativeCodeLoader: Unable to load native-hadoop library fo
r your platform... using builtin-java classes where applicable
[1, 3, 5, 7, 9]
(base) yunhe@yunhe-desktop:~$ 3-
```


- ◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡

Assignment-3 (4.0 pts.)

- DataFrames (2 pts.)
- Install PySpark (1 pt.)
- Run PySpark scripts (1 pt.)